

Apple M3 Max host (16 cores, 64 GB)

M1

VM, 3 vCPU, 6 GB

- orderer1 (Raft)
- Police peer
- CouchDB
- Police CA
- Orderer CA
- chaincode
- node-exporter

M2

VM, 3 vCPU, 6 GB

- orderer2 (Raft)
- Forensics peer
- Prosecution peer
- 2 x CouchDB
- 2 x CA
- 2 x chaincode
- node-exporter

M3

VM, 3 vCPU, 6 GB

- orderer3 (Raft)
- Court peer
- Audit peer
- 2 x CouchDB
- 2 x CA
- 2 x chaincode
- node-exporter

M4

VM, 3 vCPU, 6 GB

- DIAS backend (API)
- review store (off-chain)
- load generator
- Prometheus
- node-exporter

M5

the Mac itself (GPU)

- MLX model server
 - Qwen3-14B 4-bit
 - + V7 LoRA adapter
- one request
at a time

LLM call
(HTTP)

Docker Swarm overlay network 'diasnet' (Lima user-v2 links, TLS)